

Report

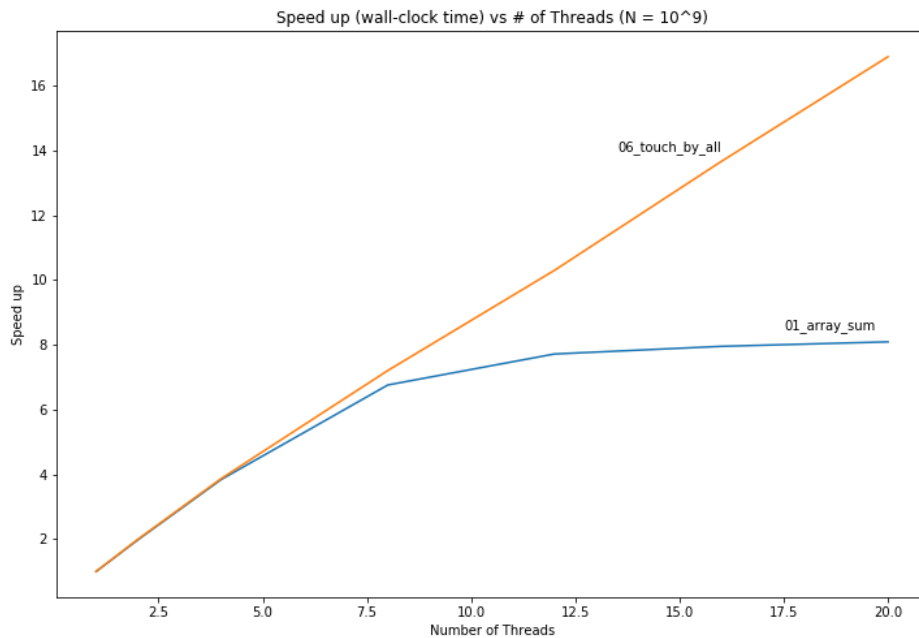
December 13, 2019

1 Exercise 0

1.1 Strong Scaling Test


1-measure the time-to-solution of the two codes in a strong-scaling test (use some meaningful value for N , like 10^9), using from 1 (using the serial version) to N_c cores on a node;

In order to test the strong scaling of codes, $N = 10^9$ chosen and from each program was run using from 1 (serial versions) to 20 threads. Both wall-clock time and elapsed time considered and tested for each step.



- 01_array_sum_output.txt speed up: [1, 1.97551078, 3.83930992, 6.76027407, 7.71710233, 7.95280571, 8.09099644]
- 06_touch_by_all_output.txt speed up: [1, 1.99257554, 3.86824436, 7.20652361, 10.30201848, 13.66230185, 16.8955089]

Above plots shows us, touch by all and touch by first scale similar up to 4 threads, after this point touch by first stopped scaling and touch by all continued to scale.



- 1 threads run has elapsed time: 6.19
- 2 threads run has elapsed time: 4.79
- 4 threads run has elapsed time: 4.11
- 8 threads run has elapsed time: 3.81
- 12 threads run has elapsed time: 3.77
- 16 threads run has elapsed time: 3.77
- 20 threads run has elapsed time: 3.78
- 01_array_sum_output.txt speed up: [1, 1.29227557, 1.50608273, 1.62467192, 1.64190981, 1.64190981, 1.63756614]
- 1 threads run has elapsed time: 6.19
- 2 threads run has elapsed time: 3.15
- 4 threads run has elapsed time: 1.62
- 8 threads run has elapsed time: 0.88
- 12 threads run has elapsed time: 0.62
- 16 threads run has elapsed time: 0.52
- 20 threads run has elapsed time: 0.48


- 06_touch_by_all_output.txt speed up: [1, 1.96507937, 3.82098765, 7.03409091, 9.98387097, 11.90384615, 12.89583333]

Above plots shows us touch by all scales better than touch by first in terms of elapsed times. Actually touch by first doesn't scale at all. On the other hand touch by all scales very good up to 16 threads but after this point, appropriately to the Amdahl's Law it stoped scaling.

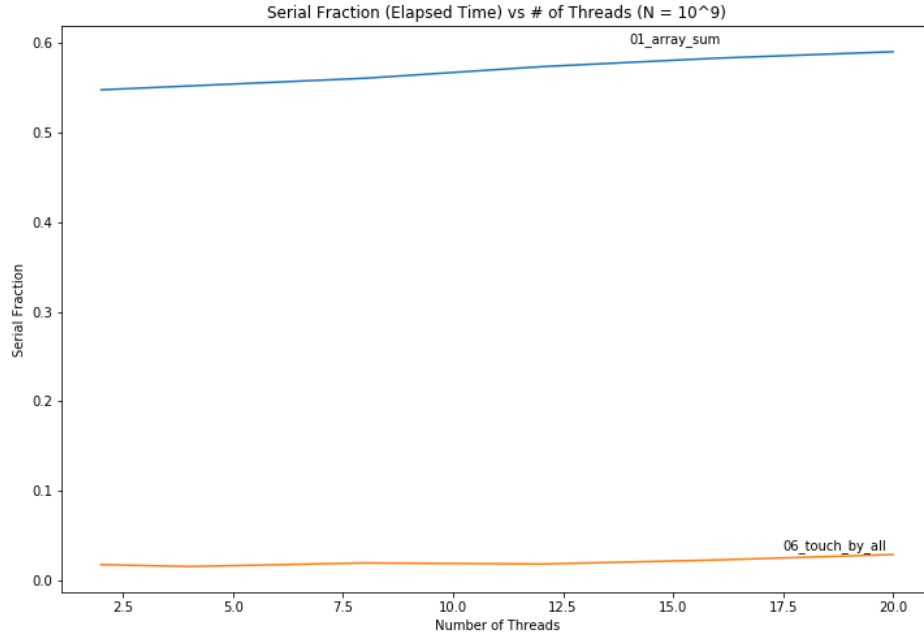
1.2 Parallel Overhead

2-measure the parallel overhead of both codes, from 2 to N_c cores on a node;

For measuring overhead two methods applied. First for both code serial fractions are calculated to understand if there is overhead or not. After that to estimate it one of the optimistic formula is used.



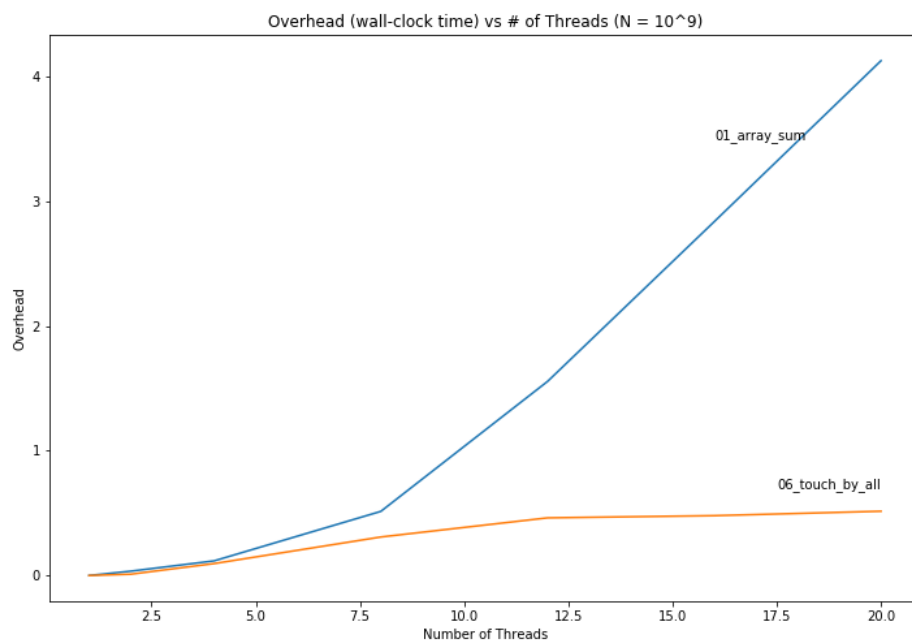
- Serial Fraction of 01_array_sum_output.txt : [0.0123964, 0.0139513, 0.02619771, 0.05045344, 0.06745791, 0.07746755]
- Serial Fraction of 06_touch_by_all_output.txt : [0.00372606, 0.01135361, 0.01572933, 0.01498366, 0.01140705, 0.00967087]




- Serial Fraction of 01_array_sum_output.txt : [0.54765751, 0.55196554, 0.56058158, 0.57350565, 0.58298331, 0.5901709]
- Serial Fraction of 06_touch_by_all_output.txt : [0.0177706, 0.01561659, 0.01961689, 0.01835806, 0.02294023, 0.02899413]

According to serial fraction calculations change (increasing serial fraction means lack of scaling is also due to the parallelization overhead) on the other hand if it is stable lack of scaling is due to the serial workload

In order to measure estimated overhead for the codes, I have used general formula, overhead function $T_o = p \times T_p - T_S$ [reference \(page 2\)](#)



- Overhead of 01_array_sum_output.txt : [0, 0.03476, 0.11736, 0.514216, 1.556208, 2.83732, 4.12722]
- Overhead of 06_touch_by_all_output.txt : [0, 0.01045, 0.095526, 0.308798, 0.46225, 0.479878, 0.51533]



- Overhead of 01_array_sum_output.txt : [0, 3.39, 10.25, 24.29, 39.05, 54.13, 69.41]
- Overhead of 06_touch_by_all_output.txt : [0, 0.11, 0.29, 0.85, 1.25, 2.13, 3.41]

In terms of overhead with the increasing number of computation units touch by first method's overhead increase too much so this situation shows us touch by all method is more efficient than touch by first method. In order to understand this difference, deeper analyze must be performed. According to this context these codes will be profiled by using perf.

1.3 Profiling Codes

3-provide any relevant metrics that explain any observed difference;

In order to specify difference between two codes, I have used perf to profile codes (collecting hardware, software events). There is no significant difference between two codes. In order to get statistically significant results data collection repeated 10 times. Results are from Ulysses (20 threads)

```
**Ulysses Computational Node** (with 20 threads)

Performance counter stats for './01_array_sum_parallel.x 1000000000' (10 runs):

    33845418144 cpu-cycles                #    3.226 GHz                ( +- 0.12% ) [80.02%]
    27367678869 instructions              #    0.81  insns per cycle    ( +- 0.05% ) [89.99%]
                                           #    0.93  stalled cycles per insn ( +- 0.06% ) [89.93%]
    180765348  cache-references           # 17.229 M/sec                ( +- 0.02% ) [89.94%]
    150130073  cache-misses                # 83.052 % of all cache refs  ( +- 0.10% ) [89.96%]
    2086147013 branch-instructions        # 198.836 M/sec               ( +- 0.83% ) [90.00%]
    119766     branch-misses              #    0.01% of all branches   ( +- 0.16% ) [90.02%]
    25395178982 stalled-cycles-frontend   # 75.03% frontend cycles idle
<not supported> stalled-cycles-backend
    10491.858119 cpu-clock                 ( +- 0.12% )
    10491.814773 task-clock                ( +- 0.12% )
<not supported> L1-dcache-loads
    380863046  L1-dcache-load-misses       #    0.00% of all L1-dcache hits ( +- 0.03% ) [90.04%]
    84430448   LLC-loads                  #    8.047 M/sec              ( +- 0.12% ) [90.05%]
    58937920   LLC-load-misses            #   69.81% of all LL-cache hits ( +- 0.11% ) [90.05%]

    3.798835260 seconds time elapsed      ( +- 0.12% )

Performance counter stats for './06_touch_by_all_parallel.x 1000000000' (10 runs):

    28103325429 cpu-cycles                #    3.066 GHz                ( +- 0.67% ) [80.00%]
    27395801411 instructions              #    0.97  insns per cycle    ( +- 0.10% ) [89.98%]
                                           #    0.72  stalled cycles per insn ( +- 0.17% ) [89.94%]
    142379842  cache-references           # 15.533 M/sec                ( +- 0.12% ) [89.95%]
    132912631  cache-misses                # 93.351 % of all cache refs  ( +- 0.23% ) [89.97%]
    2099494448 branch-instructions        # 229.042 M/sec               ( +- 0.84% ) [89.99%]
    136217     branch-misses              #    0.01% of all branches   ( +- 1.01% ) [90.01%]
    19606804162 stalled-cycles-frontend   # 69.77% frontend cycles idle
<not supported> stalled-cycles-backend
    9166.426081 cpu-clock                 ( +- 0.72% )
    9166.408297 task-clock                ( +- 0.72% )
<not supported> L1-dcache-loads
    380786971  L1-dcache-load-misses       #    0.00% of all L1-dcache hits ( +- 0.04% ) [90.03%]
    13799532   LLC-loads                  #    1.505 M/sec              ( +- 2.06% ) [90.05%]
    9623979   LLC-load-misses            #   69.74% of all LL-cache hits ( +- 2.03% ) [90.07%]

    0.525183551 seconds time elapsed      ( +- 2.25% )
```

Since there is no significant difference between two codes. Chosen events and small differences will be explained.

- Best case for cpu-cycle and instructions must be multiple instructions are executed in a single cycle so in terms of instruction per cycle touch by all policy is better than touch by first but

there is no big difference.

- In terms of cache misses results for both code is similar even touch by first is bit better than touch by all. It seems that both codes don't perform well about cache so it cause execution delays by requiring the program to fetch the data from other cache level. (perf c2c will be used to analyze this event better)
- Branch instructions and misses are more or less same for each codes and it looks efficient.
- The cycles stalled in the front-end are waste because front-end doesn't feed the back-end but for these two code percentage is more or less same
- Task clock shows time spent on the profiled task. So in this context we can say that touch by all policy is way better than touch by first since utilization of CPU (with 20 threads) in other words parallelization of touch by all (usage of threads) are more efficient than touch by first.
- To understand better, difference between two method perf c2c command used. In Ulyses perf c2c command can not be used so I used it in my local computer (with 8 threads). You can find my computers architecture and compiler info in readme file.

Trace Event Information		Trace Event Information	
Total records	: 64256	Total records	: 79561
Locked Load/Store Operations	: 666	Locked Load/Store Operations	: 1113
Load Operations	: 31716	Load Operations	: 37588
Loads - uncachable	: 0	Loads - uncachable	: 0
Loads - IO	: 0	Loads - IO	: 0
Loads - Miss	: 0	Loads - Miss	: 0
Loads - no mapping	: 1	Loads - no mapping	: 0
Load Fill Buffer Hit	: 772	Load Fill Buffer Hit	: 1743
Load L1D Hit	: 30870	Load L1D Hit	: 35243
Load L2D Hit	: 18	Load L2D Hit	: 318
Load LLC Hit	: 32	Load LLC Hit	: 260
Load Local HITM	: 0	Load Local HITM	: 79
Load Remote HITM	: 0	Load Remote HITM	: 0
Load Remote HIT	: 0	Load Remote HIT	: 0
Load Local DRAM	: 23	Load Local DRAM	: 24
Load Remote DRAM	: 0	Load Remote DRAM	: 0
Load MESI State Exclusive	: 23	Load MESI State Exclusive	: 24
Load MESI State Shared	: 0	Load MESI State Shared	: 0
Load LLC Misses	: 23	Load LLC Misses	: 24
LLC Misses to Local DRAM	: 100.0%	LLC Misses to Local DRAM	: 100.0%
LLC Misses to Remote DRAM	: 0.0%	LLC Misses to Remote DRAM	: 0.0%
LLC Misses to Remote cache (HIT)	: 0.0%	LLC Misses to Remote cache (HIT)	: 0.0%
LLC Misses to Remote cache (HITM)	: 0.0%	LLC Misses to Remote cache (HITM)	: 0.0%
Store Operations	: 32540	Store Operations	: 41973
Store - uncachable	: 0	Store - uncachable	: 0
Store - no mapping	: 1320	Store - no mapping	: 1957
Store L1D Hit	: 30635	Store L1D Hit	: 38638
Store L1D Miss	: 585	Store L1D Miss	: 1378
No Page Map Rejects	: 7899	No Page Map Rejects	: 13908
Unable to parse data source	: 0	Unable to parse data source	: 0
Global Shared Cache Line Event Information		Global Shared Cache Line Event Information	
Total Shared Cache Lines	: 0	Total Shared Cache Lines	: 35
Load Hits on shared lines	: 0	Load Hits on shared lines	: 1246
Fill Buffer Hits on shared lines	: 0	Fill Buffer Hits on shared lines	: 365
L1D hits on shared lines	: 0	L1D hits on shared lines	: 611
L2D hits on shared lines	: 0	L2D hits on shared lines	: 121
LLC hits on shared lines	: 0	LLC hits on shared lines	: 149
Locked Access on shared lines	: 0	Locked Access on shared lines	: 458
Store Hits on shared lines	: 0	Store Hits on shared lines	: 373
Store L1D hits on shared lines	: 0	Store L1D hits on shared lines	: 343
Total Merged records	: 0	Total Merged records	: 452

According to above image, touch by all (right) seems more efficient than touch by first (left) because the cache of each thread is warmed-up with the data. Cache hits are better for touch by all. There is no global shared cache line event for touch by first but touch by all has this event and naturally performs better.

2 Exercise 1

I tried to openmp-ize serial application of monte carlo pi. After few trials and modifications I got better run times than serial one. During development process generating random numbers were bit hard, first I applied standard routines but estimations of pi was terrible. That is why I changed

the way and with the help of Appendix 1 and some google search (drand48_r function requires structure) I could obtain better code. Also opening regular parallel regions (without specifying private, shared) doesn't give better results than the serial one.

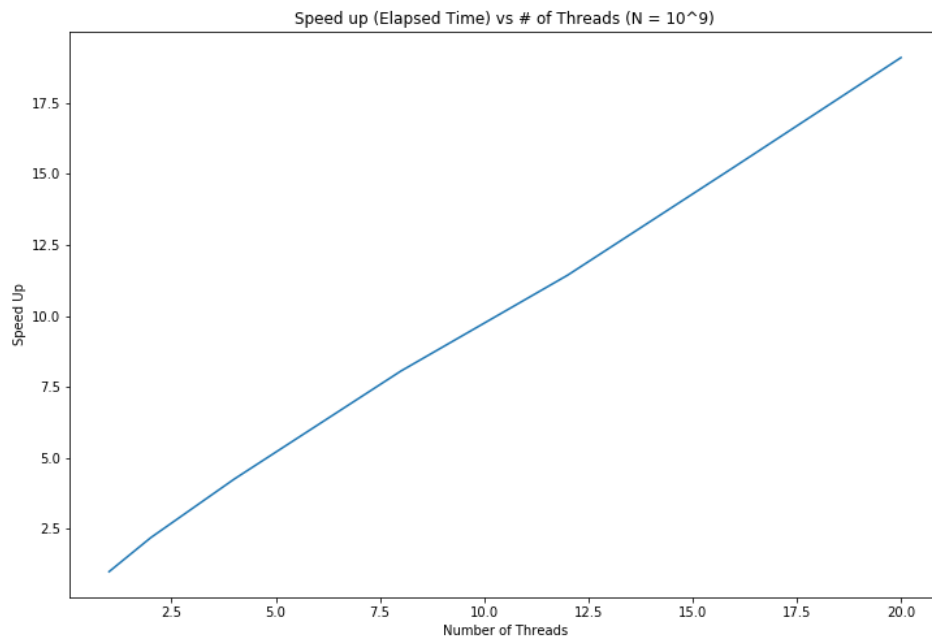
Refer the code `openmp_pi.c`

2.1 Weak and Strong Scalability

1-establish its weak and strong scalability;

2.1.1 Strong Scalability

Since elapsed time and walltime is more or less same, for this tests I'll only use elapsed time.

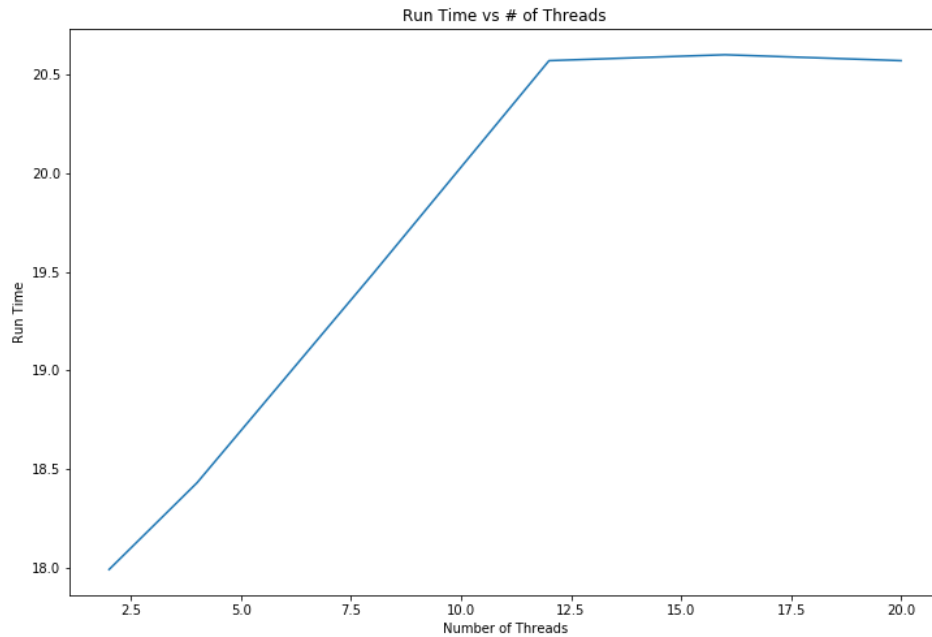


- 1 threads run has elapsed time: 19.66
- 2 threads run has elapsed time: 8.95
- 4 threads run has elapsed time: 4.62
- 8 threads run has elapsed time: 2.44
- 12 threads run has elapsed time: 1.72
- 16 threads run has elapsed time: 1.29
- 20 threads run has elapsed time: 1.03

Speed up: [1, 2.19664804, 4.25541126, 8.05737705, 11.43023256, 15.24031008, 19.08737864]

According to result of program scales linearly for $N = 10^9$.

2.1.2 Weak Scaling

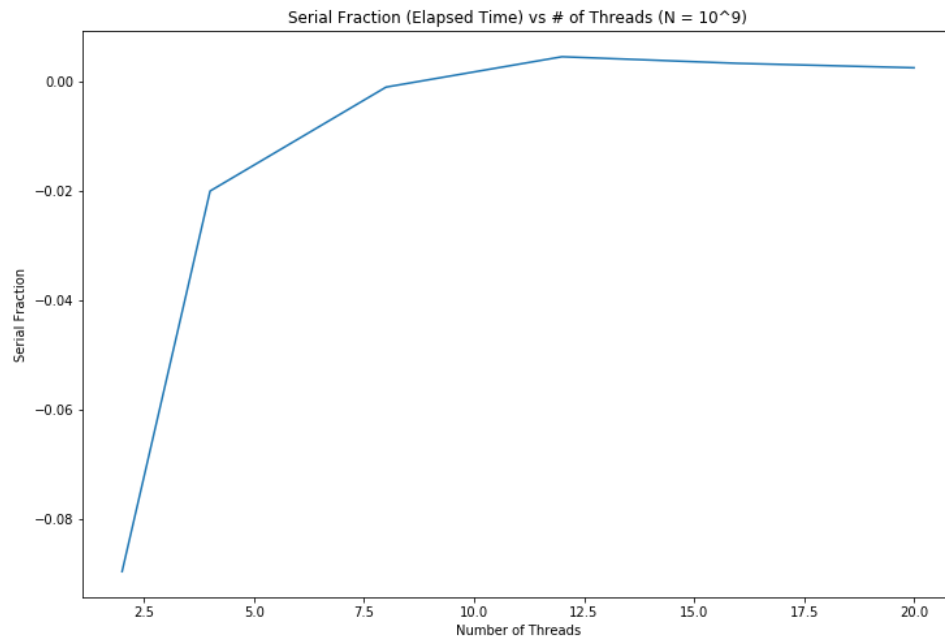


- 2 cores run has elapsed time: 17.99
- 4 cores run has elapsed time: 18.43
- 8 cores run has elapsed time: 19.49
- 12 cores run has elapsed time: 20.57
- 16 cores run has elapsed time: 20.6
- 20 cores run has elapsed time: 20.57

According to logic of weak scalability when we increase the number of cores with the N run time supposed to be same. However for this example there are small differences. There might be some room for optimization of parallelization part. But after 12 threads run time became more or less constant which is good.


2.2 Parallel Overhead

2-estimate the parallel overhead;



Serial Fraction strong_scaling_output.txt : [-0.08952187, -0.02000678, -0.00101729, 0.00453158, 0.00332316, 0.00251646]

For serial fraction, above plot shows us first there is increase (which may be due to the overhead) after that it became more or less constant so in order to estimate overhead $p \times Tp - Ts$ will be used.




Overhead for strong_scaling_output.txt : [0, -1.76, -1.18, -0.14, 0.98, 0.98, 0.94]

Similar to strong scaling results, program scales perfectly up to 16 threads but after because of overhead it slowed down. It means that there might be still room for parallel optimization especially for higher number of threads (16 and 20)

2.3 Comparing with OpenMPI

3-compare the performance of your OpenMP version and of the MPI version, in terms of time-to-solution and of parallel efficiency. Run the MPI version with N_c processes (i.e. N_c = the largest number of physical threads that you have on the node) both on the single node that you use for the OpenMP version and on multiple nodes (keeping constant the number of processes). That should allow you to understand the impact of the network and how good is the shared-memory implementation of the MPI library.

2.3.1 Single Node Comparison



- 1 cores run has elapsed time: 21.5
- 2 cores run has elapsed time: 11.79
- 4 cores run has elapsed time: 6.69
- 8 cores run has elapsed time: 4.33
- 12 cores run has elapsed time: 3.6
- 16 cores run has elapsed time: 3.18
- 20 cores run has elapsed time: 2.95

Speed up: [1, 1.8235793, 3.21375187, 4.96535797, 5.97222222, 6.76100629, 7.28813559]

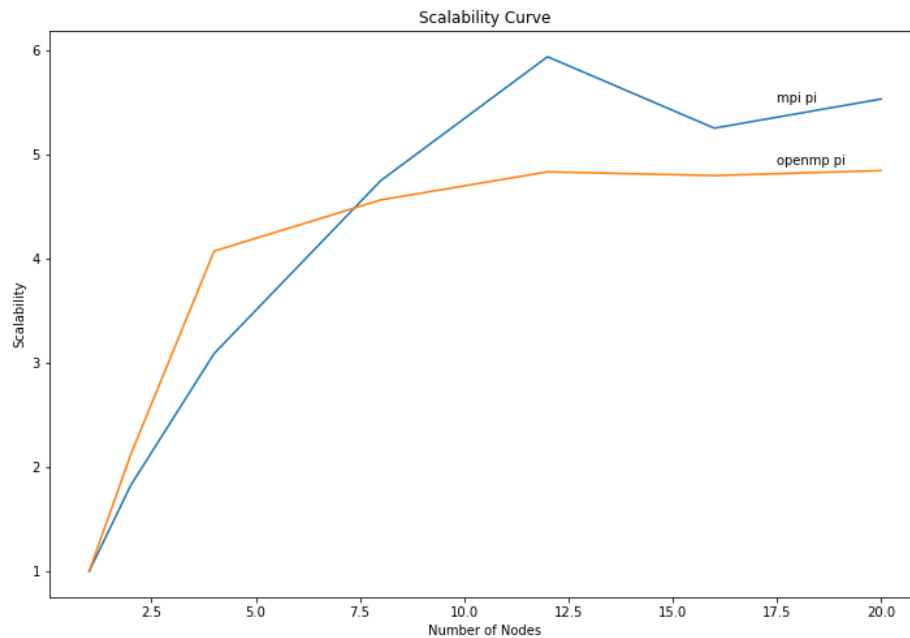
- 1 cores run has elapsed time: 19.66
- 2 cores run has elapsed time: 8.95
- 4 cores run has elapsed time: 4.62
- 8 cores run has elapsed time: 2.44
- 12 cores run has elapsed time: 1.72
- 16 cores run has elapsed time: 1.29
- 20 cores run has elapsed time: 1.03

Speed up: [1, 2.19664804, 4.25541126, 8.05737705, 11.43023256, 15.24031008, 19.08737864]

According to this comparison, in single node openmp performs better than mpi paradigm. (scalability and run time is better than mpi paradigm.) MPI program stopped scaling after 8 cores but openmp scaled linearly up to 20 threads.

2.3.2 Multiple Node Comparison

In order to understand difference between openmp and mpi paradigm same test will be done for multiple nodes with the same number of cores (I got 4 nodes and 5 cores for each node). In this way impact of network and shared memory implementation of mpi will be better understood. It doesn't make sense to run openmp program for more than one node but I run the openmp program for the same conditions to see what happens for openmp after that I will compare mpi program with single mpi performance, single node openmp performance verbally.



- 1 cores run has elapsed time: 21.86
- 2 cores run has elapsed time: 11.96
- 4 cores run has elapsed time: 7.07
- 8 cores run has elapsed time: 4.6
- 12 cores run has elapsed time: 3.68
- 16 cores run has elapsed time: 4.16
- 20 cores run has elapsed time: 3.95

Speed up: [1, 1.8277592, 3.09193777, 4.75217391, 5.94021739, 5.25480769, 5.53417722]

- 1 cores run has elapsed time: 19.68
- 2 cores run has elapsed time: 9.27
- 4 cores run has elapsed time: 4.83
- 8 cores run has elapsed time: 4.31
- 12 cores run has elapsed time: 4.07
- 16 cores run has elapsed time: 4.1
- 20 cores run has elapsed time: 4.06

Speed up: [1, 2.12297735, 4.07453416, 4.56612529, 4.83538084, 4.8, 4.84729064]

This plot shows us, in multiple nodes mpi paradigm works better. Openmp program didn't scale after 8 threads which is bigger than 5 because openmp is for shared memory. However in comparison to mpi with single node and openmp single node, it doesn't perform better than them which shows communication between nodes makes program slower.